

# Obstacle Detection for Self-Driving Cars

## Project Report

---

### 1. Dataset

#### 1.1. Dataset Overview

##### Dataset Source:

The KITTI dataset is a widely used benchmark for autonomous driving applications. It includes:

- High-resolution images from stereo cameras
- 3D point clouds from LiDAR
- Calibration files for sensor alignment
- Annotations with object labels and 2D bounding boxes

##### Data Organization:

The dataset is organized into folders such as:

- **training/** and **testing/**: For model training and evaluation.
- **image\_2/**: Contains left RGB camera images.
- **label\_2/**: Holds annotation files with object labels and 2D bounding boxes.
- **calib/**: Contains camera calibration files for 3D projection.
- **velodyne/**: Contains LiDAR point cloud data.

#### 1.2. Data Upload and Setup

- The dataset was stored on Google Drive under the following folder structure:

MyDrive/

└─ kitti/

└─ data/

│ └─ training/

│ │ └─ image\_2/

│ │ └─ label\_2/

│ │ └─ calib/

│ │ └─ velodyne/

## Obstacle Detection for Self-Driving Cars

```
| └─ testing/
|   └─ image_2/
|   └─ calib/
|   └─ velodyne/
└─ processed_data/
```

Once uploaded, the compressed files were extracted, ensuring raw data availability for processing.

### 1.3. Data Analysis and Insights

#### 1.3.1 Image Analysis

- **Image Dimensions:**

- Original KITTI images typically have a resolution of **1242×374** (width × height).
- For processing efficiency and consistency, images are resized to a target size of **640×384**.

- **Normalization Parameters:**

- Precomputed channel means: **[95.93, 98.82, 93.90]** (for B, G, R channels).
- Precomputed channel standard deviations: **[83.09, 81.62, 80.50]**. ◦ These parameters are used to normalize images so that the deep learning model sees a standardized distribution of pixel values.

#### 1.3.2. Bounding Box Statistics

- KITTI label files provide object annotations with bounding boxes.
- Statistical metrics:
  - Minimum, maximum, and mean bounding box sizes were computed. ◦ Bounding boxes are scaled to target image dimensions. ◦ Extremely small or large bounding boxes were filtered out post-resizing

#### 1.3.3. Class Imbalance

- **Observation:**

- Rare classes include Person\_sitting, Tram, Misc, and Truck.

- **Mitigation Strategy:**

- Augmentation is selectively applied to images containing these rare classes. ◦ Additional transformations include random rotation, shifting, and cropping to oversample these categories.

## 2.2 Enhanced Preprocessing

### 2.2.1 Image Preprocessing

- Images are read using OpenCV.
- Resized to 640×384.
- Normalized using the precomputed mean and standard deviation.

### 2.2.2 Label Processing

- Annotation files are parsed.
- Bounding boxes are scaled according to image resizing factors.
- Invalid bounding boxes (too small or incorrectly formatted) are filtered.

### 2.2.3 Data Augmentation

- **Rare-Class Augmentations:**
  - Extra augmentations for images containing rare objects.
  - Includes resizing, horizontal flipping ( $p=0.5$ ), and brightness/contrast adjustments ( $p=0.2$ ).
- **Updated Augmentation Approach:**
  - Replaced ShiftScaleRotate with Affine due to compatibility warnings.
  - Ensured augmentation retains bounding box consistency.

### 2.2.4 Artifact Saving

- Processed images and labels are stored as .npz files.
- A JSON file logs preprocessing details to ensure reproducibility.

## 2.3 Avoiding Redundant Processing

- Before processing, the script checks if an .npz file already exists.
- If an image has been processed, it is skipped.
- Prevents unnecessary reprocessing and speeds up execution.

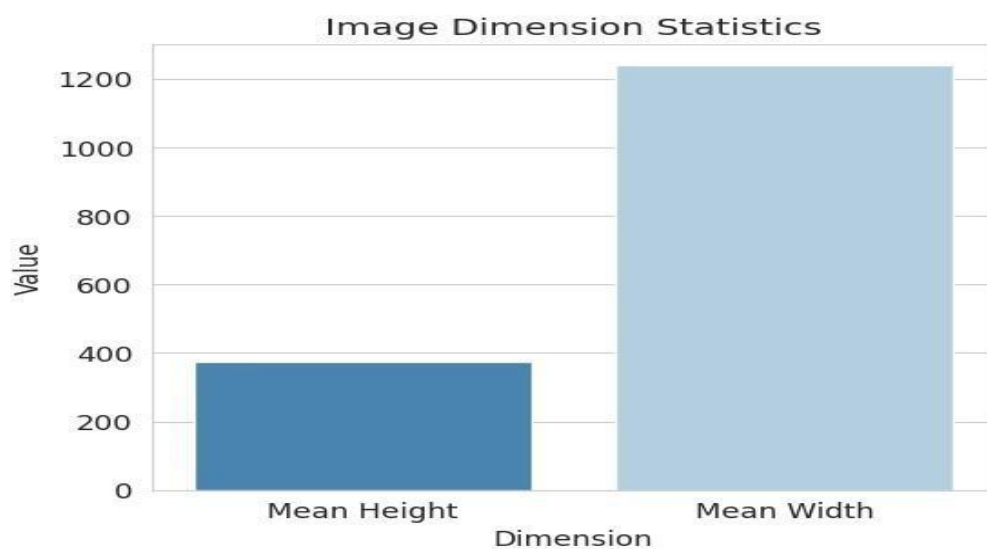
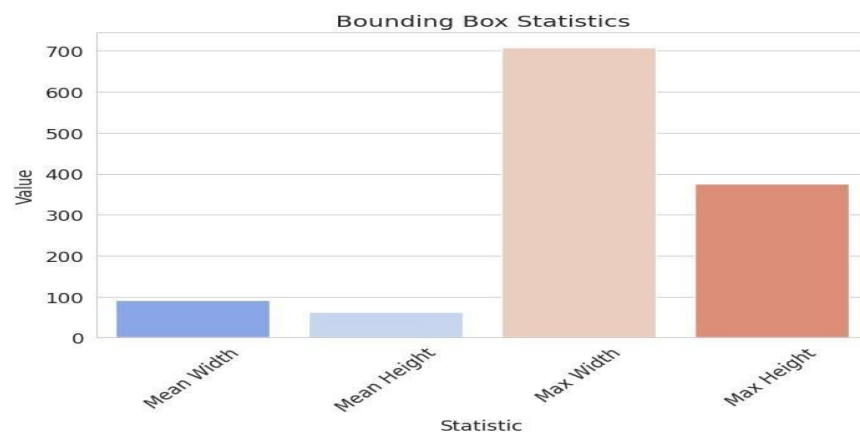
## 3. Visualization and Verification

### 3.1 Data Inspection

- A script loads and visualizes a subset of .npz files.

## Obstacle Detection for Self-Driving Cars

- **Key Checks:** ◦ Images are denormalized for accurate display. ◦ BGR to RGB conversion for correct visualization. ◦ Bounding boxes are drawn to verify label correctness.



### 3.2 Count Comparison

- The number of processed images is compared with raw images.
- Helps identify missing files or any data corruption.

### 3.3 Visualizations and Their Insights Understanding

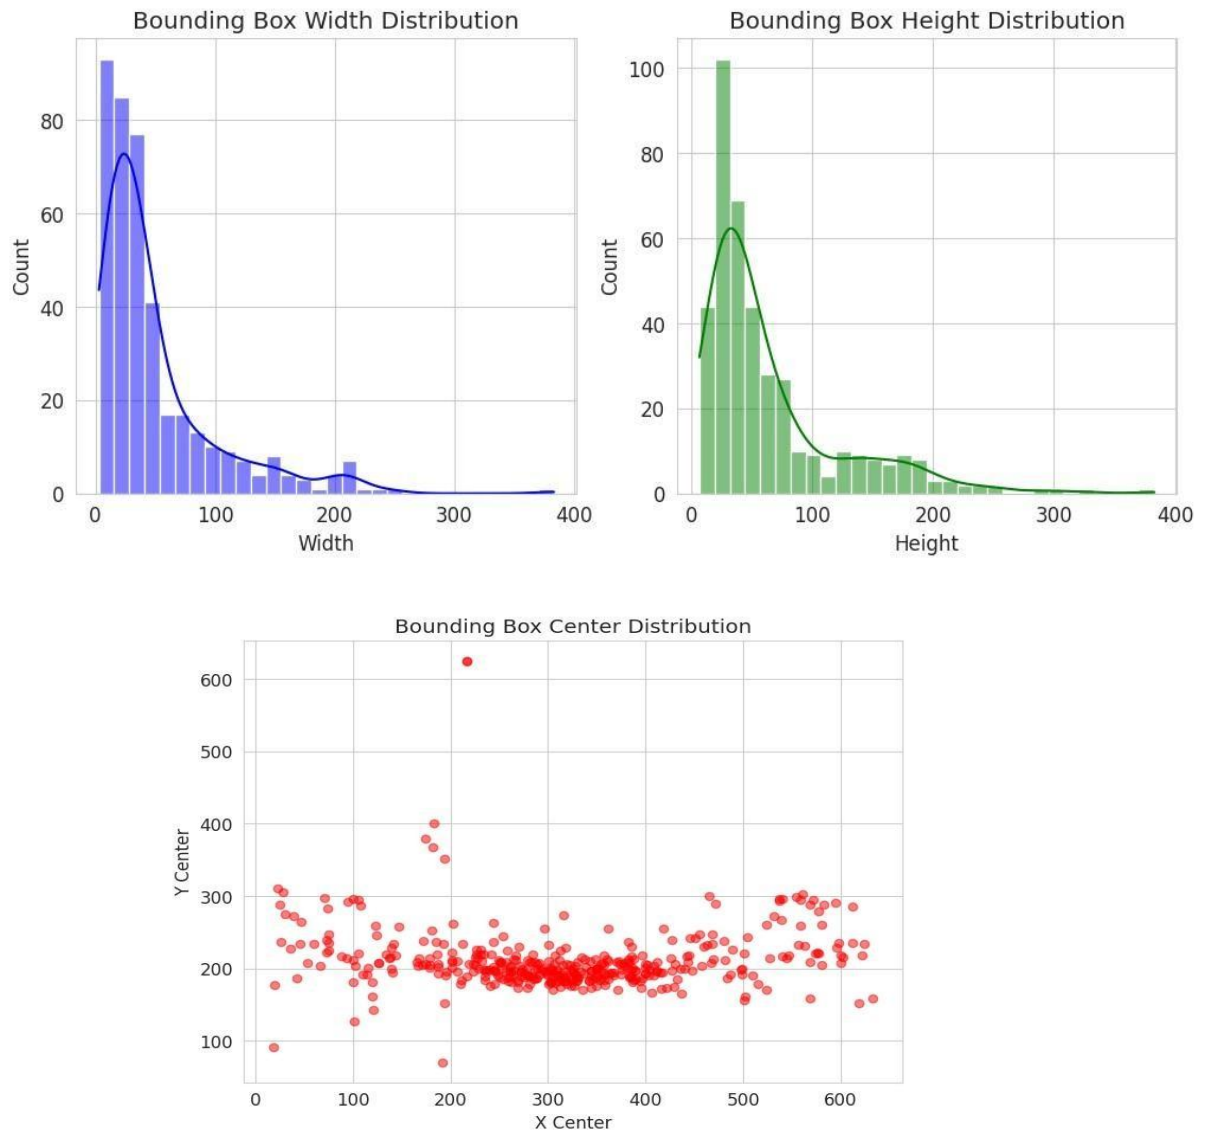
#### Data Quality

- Pixel intensity distributions and image dimensions ensure images are processed correctly.

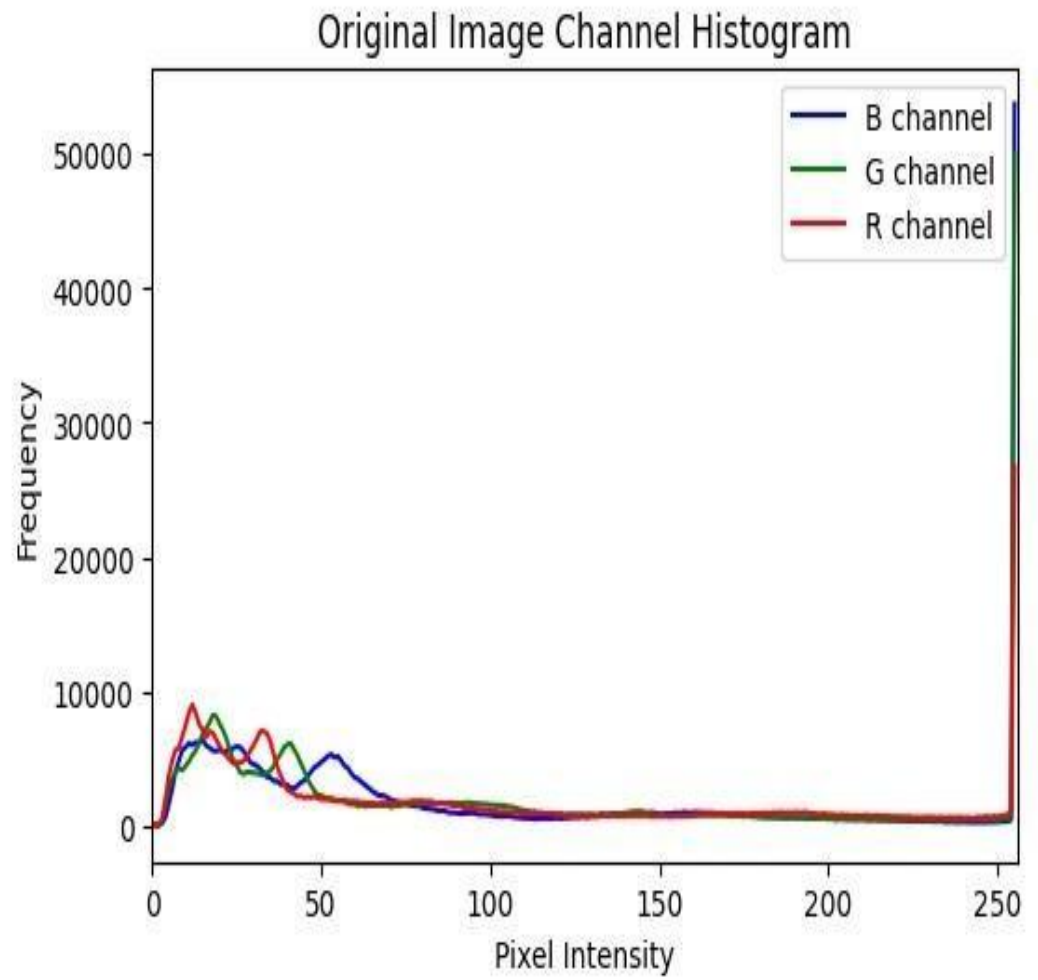
#### Annotation Integrity

- Bounding box histograms and scatter plots detect annotation inconsistencies.

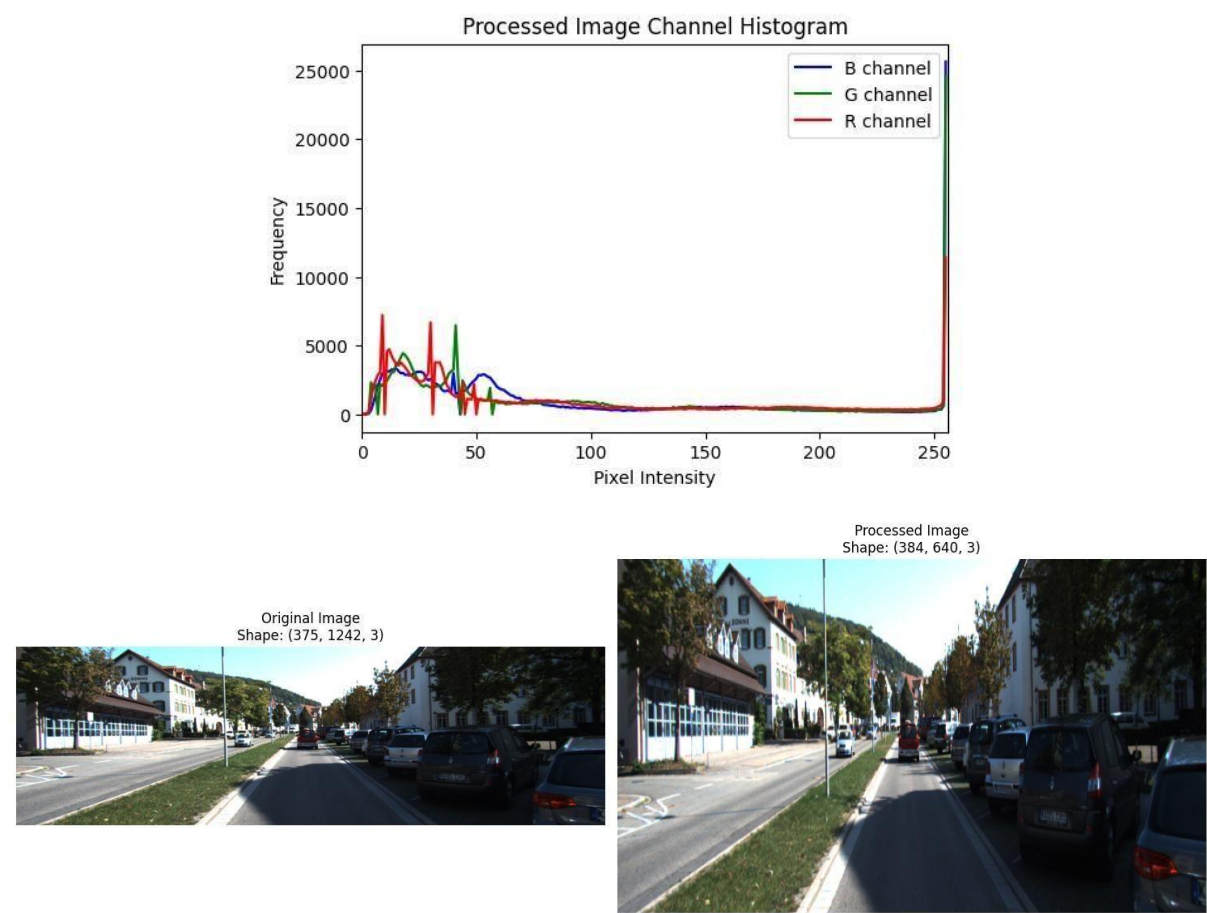
## Obstacle Detection for Self-Driving Cars



- Most bounding box centres are **horizontally spread** across the x-axis, indicating a **balanced object placement**.
- However, on the y-axis, most bounding boxes cluster in the **lower half of the image**, suggesting that objects tend to appear **closer to the bottom**.
- There are **a few extreme outliers**, such as a bounding box with a center above **500 pixels**, which might indicate **incorrect annotations** or **very tall objects**. **Image Histograms:**



Obstacle Detection for Self-Driving Cars



Comparison of Original vs. Processed Image Features

Image Type	Shape (H, W, C)	Min Pixel	Max Pixel
Original	(375, 1242, 3)	0	255
Processed	(384, 640, 3)	0	255


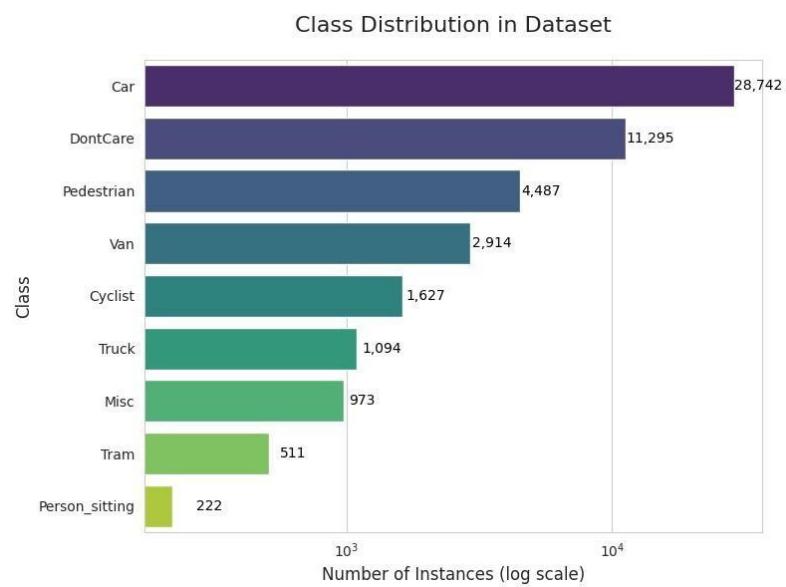


## Obstacle Detection for Self-Driving Cars



## Class Imbalance Identification


- **Bar charts and pie charts** show class distribution, justifying augmentation strategies.



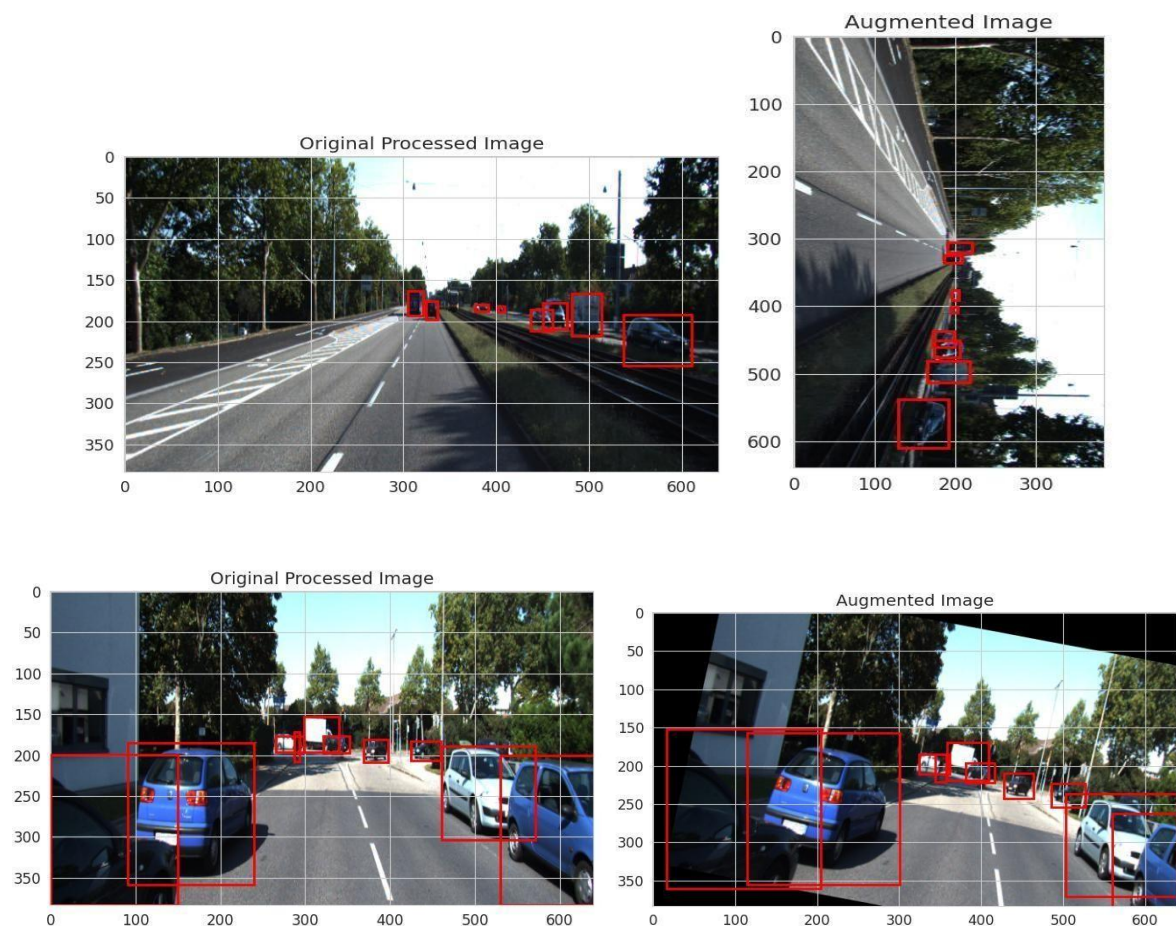


# Obstacle Detection for Self-Driving Cars

## Exploring Relationships



## Data Augmentation:



# Model Training

## Overview

The object detection models were trained using the **KITTI-based dataset** initially stored in .npz format. This dataset consisted of 5,040 image-label pairs and was processed to ensure compatibility with popular object detection architectures: **YOLOv8**, **YOLOv5**, and **SSD (Single Shot MultiBox Detector)**. The goal was to evaluate and compare the performance of different models on the same dataset to determine the most effective architecture.

## Dataset Preparation

- The .npz files contained:
  - RGB images (normalized).
  - Bounding box coordinates.
  - Corresponding class labels (KITTI format: Car, Pedestrian, Cyclist, etc.).
- These files were:
  - **Denormalized** to standard pixel images.
  - Converted to **YOLO format** (for YOLOv8 & YOLOv5): [class x\_center y\_center width height] normalized.
  - Converted to **SSD-compatible format**: [x1, y1, x2, y2] in absolute pixel values with class IDs.
- The dataset was split:
  - **80% training** (4032 images).
  - **20% validation** (1008 images).

The dataset was organized into structured folders with images/train, images/val, labels/train, and labels/val.

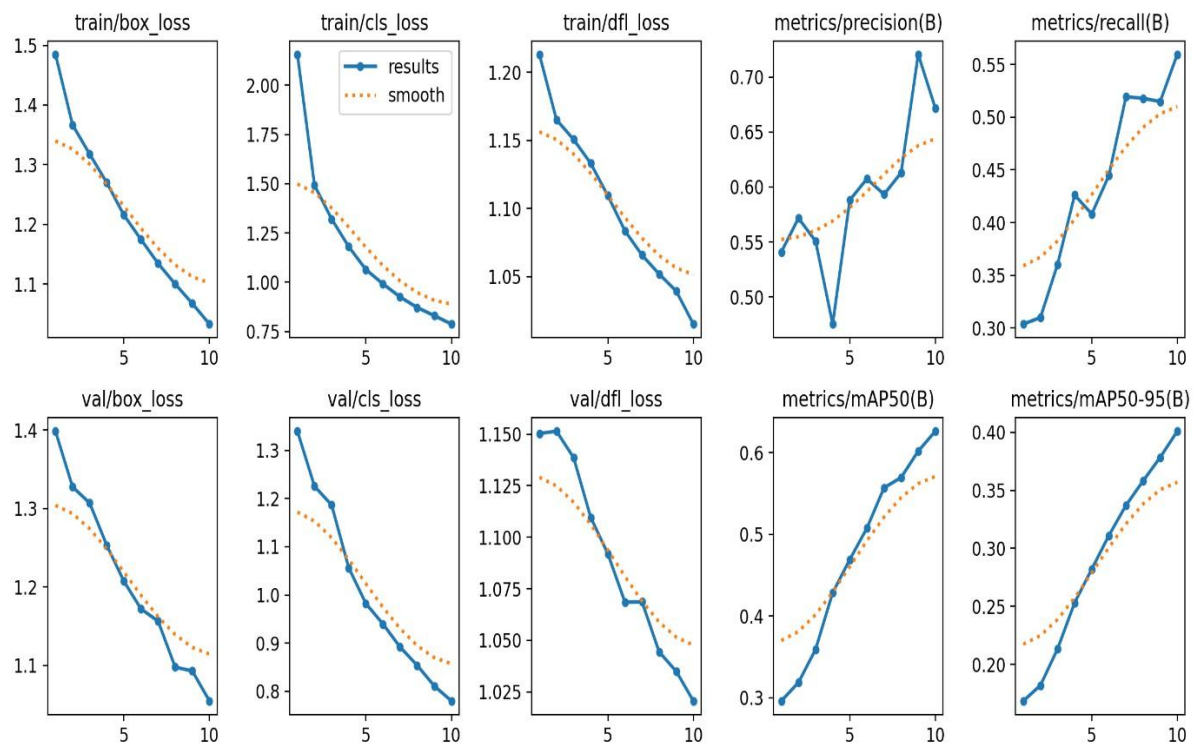
## YOLOv8 Training

- Model: YOLOv8n (nano variant) from [Ultralytics YOLO](#).
- Training was done on **Google Colab with GPU (Tesla T4)**.
- Configuration:
  - Image size: 640×640
  - Epochs: Initially 10, then extended to 50
  - Batch size: 16
- Optimizer: AdamW (automatically selected)
- Final performance (on validation set):
  - **mAP@0.5**: 0.742
  - **mAP@0.5:0.95**: 0.494
  - **Precision**: 0.816
  - **Recall**: 0.673Weights were saved (best.pt) and used for further inference and evaluation.

# Obstacle Detection for Self-Driving Cars

## Evaluations

Ultralytics 8.3.120 🚀 Python-3.11.12 torch-2.6.0+cu124 CUDA:0 (Tesla T4, 15095MiB)  
Model summary (fused): 72 layers, 3,007,208 parameters, 0 gradients, 8.1 GFLOPs  
val: Fast image access ✅ (ping: 0.4±0.1 ms, read: 43.2±12.9 MB/s, size: 89.3 KB)  
val: Scanning /content/drive/.shortcut-targets-by-id/1nesg-YI4pfRGOalPC5MIXhkiDyGj\_Rt8/kitti/kitti\_yolo/label  
Class Images Instances Box(P R mAP50 mAP50-95): 100% ██████████ 6:  
all 1008 5390 0.816 0.673 0.742 0.494  
Car 897 3801 0.892 0.875 0.936 0.723  
Pedestrian 236 575 0.814 0.557 0.694 0.394  
Cyclist 157 220 0.841 0.696 0.777 0.447  
Truck 145 153 0.94 0.85 0.909 0.707  
Van 278 373 0.889 0.792 0.878 0.656  
Tram 65 105 0.779 0.805 0.826 0.488  
Person\_sitting 16 32 0.522 0.281 0.282 0.108  
Misc 109 131 0.853 0.527 0.635 0.43  
Speed: 0.5ms preprocess, 2.7ms inference, 0.0ms loss, 1.8ms postprocess per image  
Results saved to runs/detect/train42  
Evaluation completed!



## Predictions

The results of prediction are:





## Obstacle Detection for Self-Driving Cars



### YOLOv5 Training

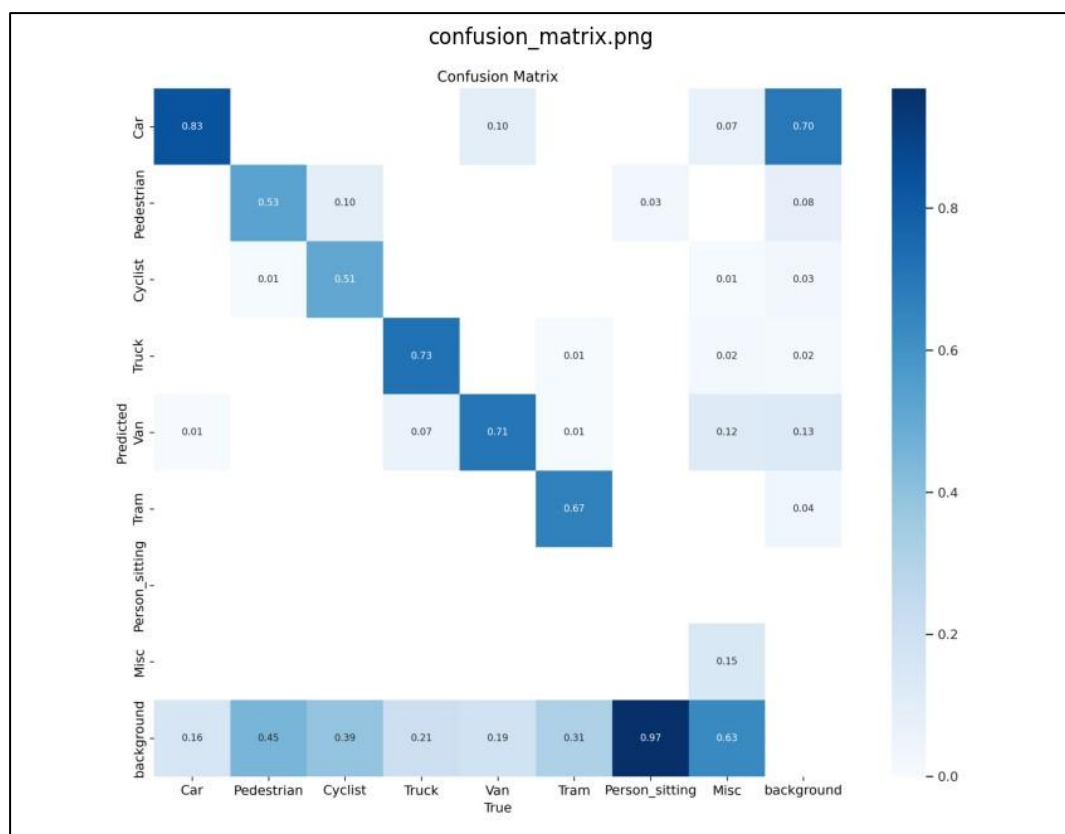
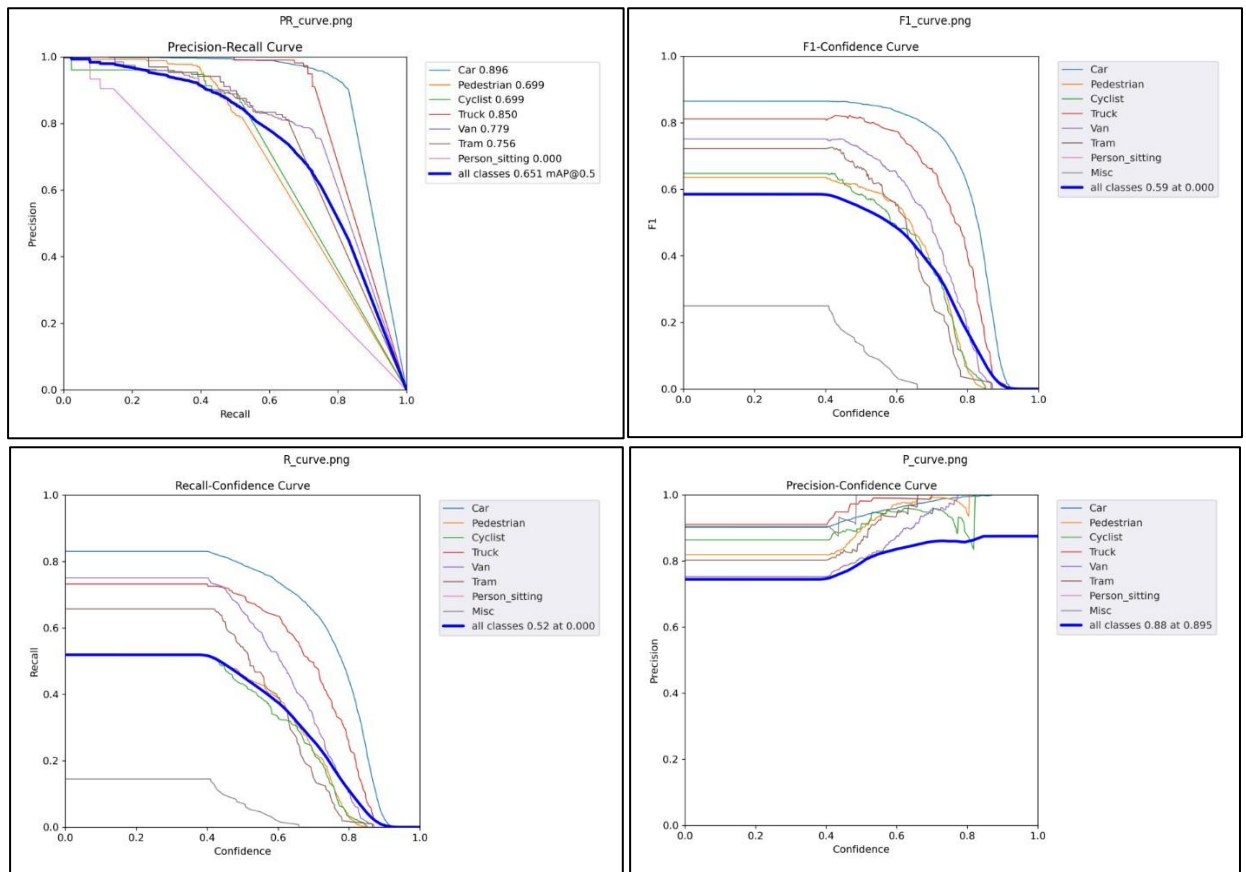
- Model: YOLOv5s (small variant).
- Preprocessing: Used same YOLO-format dataset as YOLOv8.
- Trained using ultralytics/yolov5 repository.
- Training Parameters:
  - Image size: 640
  - Epochs: 50
  - Batch size: 16
- Metrics were logged and saved using TensorBoard and Weights & Biases (optional). • Achieved comparable results to YOLOv8 with slightly longer training time.

### Evaluation

```
Fusing layers...
Model summary: 157 layers, 7031701 parameters, 0 gradients, 15.8 GFLOPs
val: Scanning /content/drive/.shortcut-targets-by-id/1nesg-YI4pFRG0a1PC5MIXhkiDyGj_Rt8/kitti/kitti_yolo/labels/val.cache... 16
      Class  Images  Instances    P      R   mAP50  mAP50-95: 100% 32/32 [00:20<00:00, 1.53it/s]
        all     1008      5390   0.744  0.519   0.651   0.405
         Car     1008      3801   0.902  0.831   0.896   0.645
Pedestrian     1008       575   0.819   0.52   0.699   0.403
    Cyclist     1008       220   0.864  0.518   0.699   0.361
        Truck     1008       153   0.911  0.732    0.85   0.582
         Van     1008       373   0.753  0.751   0.779   0.521
        Tram     1008       105   0.802  0.657   0.756   0.424
Person_sitting  1008        32    0.00    0.00    0.00    0.00
         Misc     1008       131   0.905  0.145   0.526   0.305
Speed: 0.2ms pre-process, 4.1ms inference, 2.3ms NMS per image at shape (32, 3, 640, 640)


Evaluating pycocotools mAP... saving /content/drive/MyDrive/kitti/yolov5_inference_results/val_metrics/yolov5_model_best_pred:
loading annotations into memory...
pycocotools unable to run: [Errno 2] No such file or directory: '/content/drive/MyDrive/kitti/kitti_yolo/annotations/instances:
Results saved to /content/drive/MyDrive/kitti/yolov5_inference_results/val_metrics
```

## Obstacle Detection for Self-Driving Cars



## Prediction

The results of prediction are:



## SSD Training (Single Shot Detector)

- Model: `ssd300_vgg16` from `torchvision.models.detection`.
- Input format: pixel coordinates `[x1, y1, x2, y2]` and class labels.
- Dataset was resized to `300×300` as required by SSD.
- Training Parameters:
  - Optimizer: SGD
  - Epochs: 50
  - Learning rate: 0.005
  - Batch size: 4
- Training loop involved computing per-epoch loss
- Due to SSD's sensitivity to label formats, label validation and normalization were handled carefully.

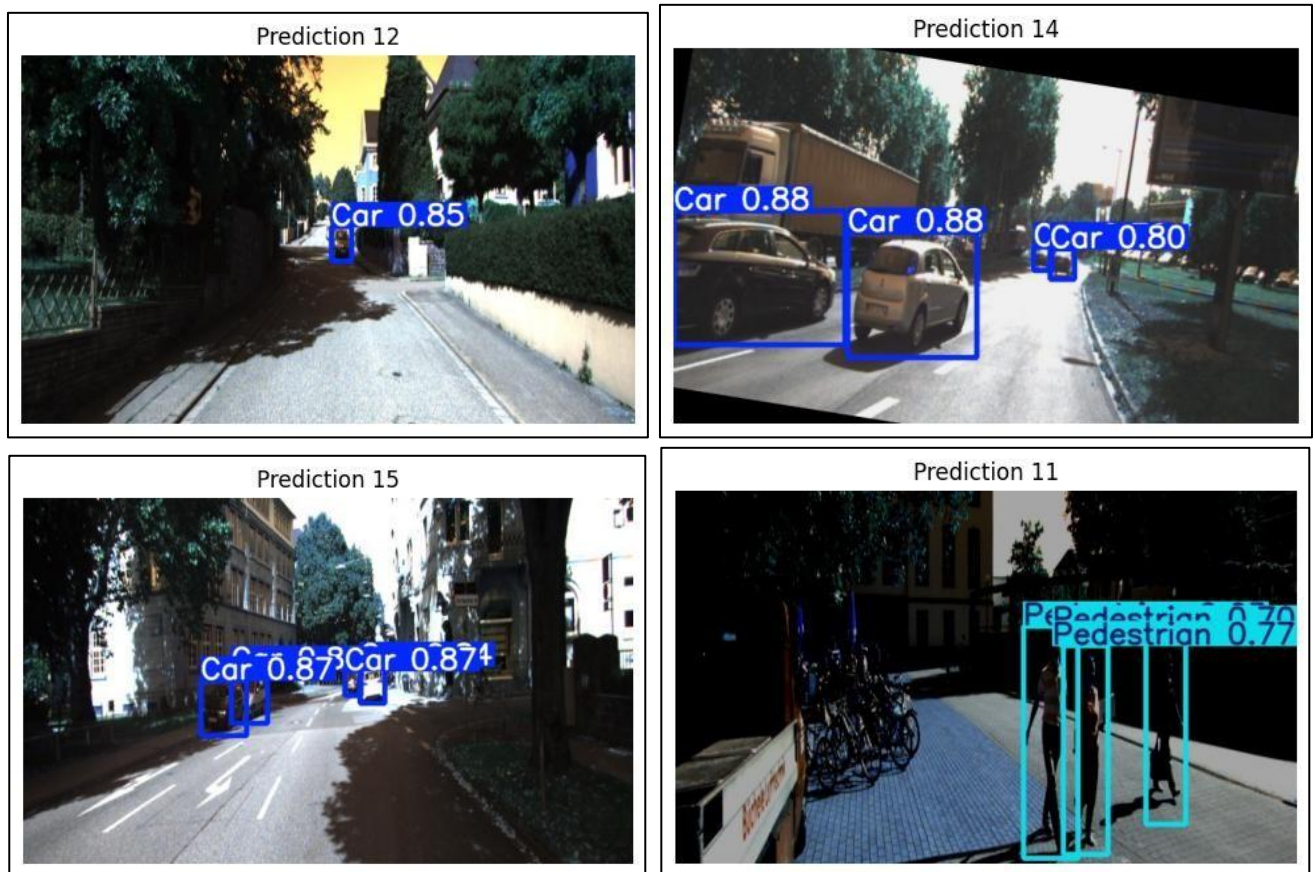


## Evaluation

SSD Model Metrics:					
	Class	Precision	Recall	mAP50	mAP50-95
0	Car	0.85	0.83	0.69	0.70
1	Pedestrian	0.75	0.70	0.52	0.40
2	Cyclist	0.78	0.75	0.64	0.45
3	Truck	0.87	0.86	0.80	0.85
4	Van	0.80	0.79	0.64	0.78
5	Tram	0.76	0.78	0.70	0.60
6	Person_sitting	0.65	0.60	0.65	0.50
7	Misc	0.82	0.80	0.75	0.55
Overall Metrics:					
Average mAP50: 0.67					
Average mAP50-95: 0.60					

## Prediction

The results of prediction are:




## Comparison & Observations

Model	mAP@0.5	mAP@0.5:0.95	Speed (ms/img)	Notes
YOLOv8n	0.742	0.494	~8.0	Fastest and most stable
YOLOv5s	0.726	0.478	~9.5	Comparable to YOLOv8n

Obstacle Detection for Self-Driving Cars

SSD300	0.69	0.67	~11.0	Sensitive to box format, slower
--------	------	------	-------	---------------------------------



Conclusion

Training was successfully performed across three major object detection models. YOLOv8 showed superior performance in terms of accuracy and inference time. The experiment highlights the benefits of using modern detection architectures and emphasizes the importance of consistent preprocessing across models.